

# Lab2\_GDB调试页表查询过程

**基础知识：**在虚拟内存访问模式下，CPU产生的虚拟地址由MMU负责转换为物理地址。MMU首先查询TLB，若命中则直接获得物理页帧号并组合成物理地址完成访问；若未命中，则根据SATP寄存器基址逐级查询页表，在获得物理页帧号后完成访问，同时将该映射存入TLB以加速后续转换。这些功能都是由**硬件**完成的。

**调试目的：**尝试观察qemu在接收到一个访存指令时是如何一步步操作的，深入对比模拟器的实现，充分理解整个地址翻译的过程。

## 步骤：

终端1：输入make debug，然后没反应，这是正常的

终端2：输入pgrep -f qemu-system-riscv64查看PID，然后sudo gdb，然后attach [刚才的PID]，然后handle SIGPIPE nostop noprint，然后设置断点break riscv\_cpu\_tlb\_fill，然后c

终端3：make gdb，然后c，此时终端2停在了断点处，然后输入 **bt** 查看栈：

```
(gdb) bt
#0 riscv_cpu_tlb_fill (cs=0x6434914b6630, address=4096, size=0,
    access_type=MMU_INST_FETCH, mmu_idx=3, probe=false, retaddr=0)
    at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:438
#1 0x000064345681afdf in tlb_fill (cpu=0x6434914b6630, addr=4096, size=0,
    access_type=MMU_INST_FETCH, mmu_idx=3, retaddr=0)
    at /mnt/d/For_OS/qemu-4.1.1/accel/tcg/cputlb.c:878
#2 0x000064345681b689 in get_page_addr_code (env=0x6434914bf040,
    addr=4096) at /mnt/d/For_OS/qemu-4.1.1/accel/tcg/cputlb.c:1032
#3 0x00006434568400ef in tb_htable_lookup (cpu=0x6434914b6630, pc=4096,
    cs_base=0, flags=3, cf_mask=4278714368)
    at /mnt/d/For_OS/qemu-4.1.1/accel/tcg/cpu-exec.c:339
#4 0x000064345683f62d in tb_lookup__cpu_state (cpu=0x6434914b6630,
    pc=0x79cbd5ee2928, cs_base=0x79cbd5ee2920, flags=0x79cbd5ee291c,
    cf_mask=4278714368)
    at /mnt/d/For_OS/qemu-4.1.1/include/exec/tb-lookup.h:43
#5 0x00006434568403c2 in tb_find (cpu=0x6434914b6630, last_tb=0x0,
    tb_exit=0, cf_mask=524288)
    at /mnt/d/For_OS/qemu-4.1.1/accel/tcg/cpu-exec.c:406
#6 0x0000643456840ce9 in cpu_exec (cpu=0x6434914b6630)
    at /mnt/d/For_OS/qemu-4.1.1/accel/tcg/cpu-exec.c:731
#7 0x00006434567f2f3f in tcg_cpu_exec (cpu=0x6434914b6630)
    at /mnt/d/For_OS/qemu-4.1.1/cpus.c:1435
#8 0x00006434567f37f8 in qemu_tcg_cpu_thread_fn (arg=0x6434914b6630)
    at /mnt/d/For_OS/qemu-4.1.1/cpus.c:1743
#9 0x0000643456c768d7 in qemu_thread_start (args=0x6434914cccc0)
    at util/qemu-thread-posix.c:502
#10 0x000079cbd689caa4 in start_thread (arg=<optimized out>)
    at ./nptl/pthread_create.c:447
#11 0x000079cbd6929c6c in clone3 ()
    at ../sysdeps/unix/sysv/linux/x86_64/clone3.S:78
```

上述调用栈展示了QEMU模拟RISC-V CPU时**TLB未命中**的处理流程：当CPU执行到reset vector地址 `0x1000` (4096) 进行指令取指时，TLB中没有该地址的映射，于是触发TLB重填。调用从 `cpu_exec` (CPU主执行循环) 开始，通过 `tb_find` 查找翻译块，发现需要翻译 `0x1000` 地址的代码。`tb_htable_lookup` 调用 `get_page_addr_code` 获取物理地址，但发现页表未缓存，于是调用 `tlb_fill` 进行TLB填充。最终进入 `riscv_cpu_tlb_fill` 函数，该函数会调用 `get_physical_address` 进行实际的页表遍历，将虚拟地址 `0x1000` 转换为物理地址，并填充到TLB中供后续使用。

整个过程体现了**指令执行 → TLB未命中 → 页表遍历 → TLB填充**的完整MMU工作流程。

---

然后输入 `(gdb) list`，我们可以看到 `riscv_cpu_tlb_fill()` 的开头部分：

```
(gdb) list
433     #endif
434
435     bool riscv_cpu_tlb_fill(CPUState *cs, vaddr address, int size,
436                               MMUAccessType access_type, int mmu_idx,
437                               bool probe, uintptr_t retaddr)
438     {
439     #ifndef CONFIG_USER_ONLY
440         RISCVCPU *cpu = RISCV_CPU(cs);
441         CPURISCVState *env = &cpu->env;
442         hwaddr pa = 0;
```

然后输入 (gdb) print address 查看访存地址，输入 (gdb) print access\_type 查看访存类型，输入 (gdb) print mmu\_idx 查看MMU模式：

```
(gdb) p address
$1 = 4096
(gdb) p access_type
$2 = MMU_INST_FETCH
(gdb) p mmu_idx
$3 = 3
```

可以看到，访存地址为 \$1 = 4096，也就是 0x1000。访问类型为 MMU\_INST\_FETCH，说明是执行指令时触发翻译，这是正常的。MMU模式为3，表示当前为M模式。

接下来在终端2输入多个 (gdb) step，直至到达函数 get\_physical\_address()：

```
(gdb)
get_physical_address (env=0x6434914bf040, physical=0x79cbd5ee2720,
    prot=0x79cbd5ee2714, addr=4096, access_type=2, mmu_idx=3)
  at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:158
158     {
```

我们去对应的路径 /mnt/d/For\_OS/qemu-4.1.1/target/riscv/cpu\_helper.c 中看看这个函数的qemu源码：

#### 代码块

```
1 static int get_physical_address(CPURISCVState *env, hwaddr *physical,
2                                 int *prot, target_ulong addr,
3                                 int access_type, int mmu_idx)
4 {
5     /* NOTE: the env->pc value visible here will not be
6      * correct, but the value visible to the exception handler
7      * (riscv_cpu_do_interrupt) is correct */
8
9     int mode = mmu_idx;
10
11    if (mode == PRV_M && access_type != MMU_INST_FETCH) {
```

```

12     if (get_field(env->mstatus, MSTATUS_MPRV)) {
13         mode = get_field(env->mstatus, MSTATUS_MPP);
14     }
15 }
16
17 if (mode == PRV_M || !riscv_feature(env, RISCV_FEATURE_MMU)) {
18     *physical = addr;
19     *prot = PAGE_READ | PAGE_WRITE | PAGE_EXEC;
20     return TRANSLATE_SUCCESS;
21 }
22
23 *prot = 0;
24
25 target_ulong base;
26 int levels, ptidxbits, ptesize, vm, sum;
27 int mxr = get_field(env->mstatus, MSTATUS_MXR);
28
29 if (env->priv_ver >= PRIV_VERSION_1_10_0) {
30     base = get_field(env->satp, SATP_PPN) << PGSHIFT;
31     sum = get_field(env->mstatus, MSTATUS_SUM);
32     vm = get_field(env->satp, SATP_MODE);
33     switch (vm) {
34     case VM_1_10_SV32:
35         levels = 2; ptidxbits = 10; ptesize = 4; break;
36     case VM_1_10_SV39:
37         levels = 3; ptidxbits = 9; ptesize = 8; break;
38     case VM_1_10_SV48:
39         levels = 4; ptidxbits = 9; ptesize = 8; break;
40     case VM_1_10_SV57:
41         levels = 5; ptidxbits = 9; ptesize = 8; break;
42     case VM_1_10_MBARE:
43         *physical = addr;
44         *prot = PAGE_READ | PAGE_WRITE | PAGE_EXEC;
45         return TRANSLATE_SUCCESS;
46     default:
47         g_assert_not_reached();
48     }
49 } else {
50     base = env->sptbr << PGSHIFT;
51     sum = !get_field(env->mstatus, MSTATUS_PUM);
52     vm = get_field(env->mstatus, MSTATUS_VM);
53     switch (vm) {
54     case VM_1_09_SV32:
55         levels = 2; ptidxbits = 10; ptesize = 4; break;
56     case VM_1_09_SV39:
57         levels = 3; ptidxbits = 9; ptesize = 8; break;
58     case VM_1_09_SV48:

```

```

59         levels = 4; ptidxbits = 9; ptesize = 8; break;
60     case VM_1_09_MBARE:
61         *physical = addr;
62         *prot = PAGE_READ | PAGE_WRITE | PAGE_EXEC;
63         return TRANSLATE_SUCCESS;
64     default:
65         g_assert_not_reached();
66     }
67 }
68
69 CPUState *cs = env_cpu(env);
70 int va_bits = PGSHIFT + levels * ptidxbits;
71 target_ulong mask = (1L << (TARGET_LONG_BITS - (va_bits - 1))) - 1;
72 target_ulong masked_msbs = (addr >> (va_bits - 1)) & mask;
73 if (masked_msbs != 0 && masked_msbs != mask) {
74     return TRANSLATE_FAIL;
75 }
76
77 int ptshift = (levels - 1) * ptidxbits;
78 int i;
79
80 #if !TCG_OVERSIZED_GUEST
81 restart:
82 #endif
83     for (i = 0; i < levels; i++, ptshift -= ptidxbits) {
84         target_ulong idx = (addr >> (PGSHIFT + ptshift)) &
85                         ((1 << ptidxbits) - 1);
86
87         /* check that physical address of PTE is legal */
88         target_ulong pte_addr = base + idx * ptesize;
89
90         if (riscv_feature(env, RISCV_FEATURE_PMP) &&
91             !pmp_hart_has_privs(env, pte_addr, sizeof(target_ulong),
92             1 << MMU_DATA_LOAD, PRV_S)) {
93             return TRANSLATE_PMP_FAIL;
94         }
95 #if defined(TARGET_RISCV32)
96         target_ulong pte = ldl_phys(cs->as, pte_addr);
97 #elif defined(TARGET_RISCV64)
98         target_ulong pte = ldq_phys(cs->as, pte_addr);
99 #endif
100        target_ulong ppn = pte >> PTE_PPN_SHIFT;
101
102        if (!(pte & PTE_V)) {
103            /* Invalid PTE */
104            return TRANSLATE_FAIL;
105        } else if (!(pte & (PTE_R | PTE_W | PTE_X))) {

```

```

106     /* Inner PTE, continue walking */
107     base = ppn << PGSHIFT;
108 } else if ((pte & (PTE_R | PTE_W | PTE_X)) == PTE_W) {
109     /* Reserved leaf PTE flags: PTE_W */
110     return TRANSLATE_FAIL;
111 } else if ((pte & (PTE_R | PTE_W | PTE_X)) == (PTE_W | PTE_X)) {
112     /* Reserved leaf PTE flags: PTE_W + PTE_X */
113     return TRANSLATE_FAIL;
114 } else if ((pte & PTE_U) && ((mode != PRV_U) &&
115             (!sum || access_type == MMU_INST_FETCH))) {
116     /* User PTE flags when not U mode and mstatus.SUM is not set,
117      or the access type is an instruction fetch */
118     return TRANSLATE_FAIL;
119 } else if (!(pte & PTE_U) && (mode != PRV_S)) {
120     /* Supervisor PTE flags when not S mode */
121     return TRANSLATE_FAIL;
122 } else if (ppn & ((1ULL << ptshift) - 1)) {
123     /* Misaligned PPN */
124     return TRANSLATE_FAIL;
125 } else if (access_type == MMU_DATA_LOAD && !((pte & PTE_R) ||
126             ((pte & PTE_X) && mxr))) {
127     /* Read access check failed */
128     return TRANSLATE_FAIL;
129 } else if (access_type == MMU_DATA_STORE && !(pte & PTE_W)) {
130     /* Write access check failed */
131     return TRANSLATE_FAIL;
132 } else if (access_type == MMU_INST_FETCH && !(pte & PTE_X)) {
133     /* Fetch access check failed */
134     return TRANSLATE_FAIL;
135 } else {
136     /* if necessary, set accessed and dirty bits. */
137     target_ulong updated_pte = pte | PTE_A |
138         (access_type == MMU_DATA_STORE ? PTE_D : 0);
139
140     /* Page table updates need to be atomic with MTTCG enabled */
141     if (updated_pte != pte) {
142         /*
143          * - if accessed or dirty bits need updating, and the PTE is
144          *   in RAM, then we do so atomically with a compare and swap.
145          * - if the PTE is in IO space or ROM, then it can't be updated
146          *   and we return TRANSLATE_FAIL.
147          * - if the PTE changed by the time we went to update it, then
148          *   it is no longer valid and we must re-walk the page table.
149          */
150     MemoryRegion *mr;
151     hwaddr l = sizeof(target_ulong), addr1;
152     mr = address_space_translate(cs->as, pte_addr,

```

```

153             &addr1, &l, false, MEMTXATTRS_UNSPECIFIED);
154         if (memory_region_is_ram(mr)) {
155             target_ulong *pte_pa =
156                 qemu_map_ram_ptr(mr->ram_block, addr1);
157 #if TCG_OVERSIZED_GUEST
158             /* MTTCG is not enabled on oversized TCG guests so
159              * page table updates do not need to be atomic */
160             *pte_pa = pte = updated_pte;
161 #else
162             target_ulong old_pte =
163                 atomic_cmpxchg(pte_pa, pte, updated_pte);
164             if (old_pte != pte) {
165                 goto restart;
166             } else {
167                 pte = updated_pte;
168             }
169 #endif
170         } else {
171             /* misconfigured PTE in ROM (AD bits are not preset) or
172              * PTE is in IO space and can't be updated atomically */
173             return TRANSLATE_FAIL;
174         }
175     }
176
177     /* for superpage mappings, make a fake leaf PTE for the TLB's
178      benefit. */
179     target_ulong vpn = addr >> PGSHIFT;
180     *physical = (ppn | (vpn & ((1L << ptshift) - 1))) << PGSHIFT;
181
182     /* set permissions on the TLB entry */
183     if ((pte & PTE_R) || ((pte & PTE_X) && mxr)) {
184         *prot |= PAGE_READ;
185     }
186     if ((pte & PTE_X)) {
187         *prot |= PAGE_EXEC;
188     }
189     /* add write permission on stores or if the page is already dirty,
190      so that we TLB miss on later writes to update the dirty bit */
191     if ((pte & PTE_W) &&
192             (access_type == MMU_DATA_STORE || (pte & PTE_D))) {
193         *prot |= PAGE_WRITE;
194     }
195     return TRANSLATE_SUCCESS;
196 }
197 }
198 return TRANSLATE_FAIL;
199 }
```

这是QEMU中RISC-V架构的页表遍历核心函数，负责将虚拟地址转换为物理地址，是MMU（内存管理单元）的核心实现。

### 函数参数：

- `env` : CPU状态寄存器环境
- `physical` : 输出参数，存放转换后的物理地址
- `prot` : 输出参数，存放页面访问权限
- `addr` : 输入的虚拟地址
- `access_type` : 访问类型（取指/读/写）
- `mmu_idx` : MMU索引，对应CPU特权级

### 执行流程：

1. 模式检查与快速路径：M模式或无MMU则直接映射
2. 页表配置解析：读取sapt寄存器
3. 地址有效性检查：检查虚拟地址的高位是否符号扩展正确
4. 多级页表遍历（核心）：SV39三级遍历循环 for循环
5. PTE权限检查：
  - `PTE_V` : 有效位（必须为1）
  - `PTE_R/W/X` : 读/写/执行权限
  - `PTE_U` : 用户页标志
  - `PTE_A/D` : 访问/脏位
6. 物理地址生成：将PTE中的PPN与虚拟地址的低位组合
7. 权限设置：根据PTE权限设置 `prot` 参数
  - `PAGE_READ` : 可读
  - `PAGE_WRITE` : 可写
  - `PAGE_EXEC` : 可执行

然后我们尝试遍历三级页表：

```
(gdb) print env->priv  
$4 = 3  
(gdb) print riscv_feature(env, RISCV_FEATURE_MMU)  
$5 = true
```

代码块

```
1 (gdb) print env->priv
2 $4 = 3 #仍然是M模式,
3 (gdb) print riscv_feature(env, RISCV_FEATURE_MMU)
4 $5 = true #虽然是true, 但因为是M模式, 所以不会遍历页表
```

所以要等待S或U模式的访问，我们设置条件断点：

```
(gdb) delete 1
(gdb) break get_physical_address if env->priv != 3
Breakpoint 2 at 0x6434568d3fc2: file /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c, line 158.
(gdb) c
Continuing.

Thread 2 "qemu-system-ris" hit Breakpoint 2, get_physical_address (
    env=0x6434914bf040, physical=0x79cbd5ee2720, prot=0x79cbd5ee2714,
    addr=2149580800, access_type=2, mmu_idx=1)
    at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:158
158 {
(gdb) print env->priv
$6 = 1
(gdb) print riscv_feature(env, RISCV_FEATURE_MMU)
$7 = true
(gdb) print /x addr
$8 = 0x80200000
(gdb) print /x env->satp
$9 = 0x0
(gdb) print /x env->satp & 0xFF
$10 = 0x0
```

### 代码块

```
1 (gdb) delete 1 #删除当前断点
2 (gdb) break get_physical_address if env->priv != 3 #设置条件断点: 只在非M模式时触发
3 Breakpoint 2 at 0x6434568d3fc2: file /mnt/d/For_OS/qemu-
4.1.1/target/riscv/cpu_helper.c, line 158.
4 (gdb) c #继续执行
5 Continuing.

6

7 Thread 2 "qemu-system-ris" hit Breakpoint 2, get_physical_address (
8     env=0x6434914bf040, physical=0x79cbd5ee2720, prot=0x79cbd5ee2714,
9     addr=2149580800, access_type=2, mmu_idx=1)
10    at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:158
11 158 {
12 (gdb) print env->priv
13 $6 = 1 #现在切换到S模式了!
14 (gdb) print riscv_feature(env, RISCV_FEATURE_MMU)
15 $7 = true
16
17 (gdb) print /x addr
```

```

18 $8 = 0x80200000
19 (gdb) print /x env->satp
20 $9 = 0x0 # satp = 0, 表示分页还没有启用
21 (gdb) print /x env->satp & 0xFF
22 $10 = 0x0

```

我们继续设置断点：

```

(gdb) delete
Delete all breakpoints, watchpoints, tracepoints, and catchpoints? (y or n)
y
(gdb) break get_physical_address if env->satp != 0
Breakpoint 3 at 0x6434568d3fc2: file /mnt/d/For_OS/qemu-4.1.1/target/riscv/c
pu_helper.c, line 158.
(gdb) c
Continuing.

Thread 2 "qemu-system-ris" hit Breakpoint 3, get_physical_address (
    env=0x6434914bf040, physical=0x79cbd5ee21f0, prot=0x79cbd5ee21e4,
    addr=18446744072637907160, access_type=2, mmu_idx=1)
    at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:158
158 {
(gdb) print env->priv
$11 = 1
(gdb) print /x env->satp
$12 = 0x80000000000080206
(gdb) print /x env->satp & 0xFF
$13 = 0x6
(gdb) print /x addr
$14 = 0xffffffffc02000d8

```

### 代码块

```

1 (gdb) delete
2 Delete all breakpoints, watchpoints, tracepoints, and catchpoints? (y or n) y
3 (gdb) break get_physical_address if env->satp != 0 #继续设置条件断点
4 Breakpoint 3 at 0x6434568d3fc2: file /mnt/d/For_OS/qemu-
    4.1.1/target/riscv/cpu_helper.c, line 158.
5 (gdb) c
6 Continuing.

7

8 Thread 2 "qemu-system-ris" hit Breakpoint 3, get_physical_address (
9     env=0x6434914bf040, physical=0x79cbd5ee21f0, prot=0x79cbd5ee21e4,
10    addr=18446744072637907160, access_type=2, mmu_idx=1)
11    at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:158
12 158 {
13 (gdb) print env->priv
14 $11 = 1 # S模式✓
15 (gdb) print /x env->satp
16 $12 = 0x80000000000080206 # satp 不为0, 表示分页已启用✓

```

```
17 (gdb) print /x env->satp & 0xFF
18 $13 = 0x6
19 (gdb) print /x addr
20 $14 = 0xfffffffffc02000d8 # 高地址, 是用户空间地址✓
```

现在条件都满足了，可以开始调试了，输入下述命令：

```
(gdb) delete 3
(gdb) break cpu_helper.c:237
Breakpoint 4 at 0x6434568d440d: file /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c, line 237.
(gdb) c
Continuing.

Thread 2 "qemu-system-ris" hit Breakpoint 4, get_physical_address (
    env=0x6434914bf040, physical=0x79cbd5ee21f0, prot=0x79cbd5ee21e4,
    addr=18446744072637907160, access_type=2, mmu_idx=1)
    at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:237
237         for (i = 0; i < levels; i++, ptshift -= ptidxbits) {
(gdb) print i
$15 = -705808704
(gdb) print ptshift
$16 = 18
(gdb) print levels
$17 = 3
(gdb) print ptidxbits
$18 = 9
(gdb) print /x base
$19 = 0x80206000
```

## 代码块

```
1 (gdb) delete 3 #删除原有断点
2 (gdb) break cpu_helper.c:237 #设置新断点: get_physical_address()的for循环处
3 Breakpoint 4 at 0x6434568d440d: file /mnt/d/For_OS/qemu-
   4.1.1/target/riscv/cpu_helper.c, line 237.
4 (gdb) c
5 Continuing.

6
7 Thread 2 "qemu-system-ris" hit Breakpoint 4, get_physical_address (
8     env=0x6434914bf040, physical=0x79cbd5ee21f0, prot=0x79cbd5ee21e4,
9     addr=18446744072637907160, access_type=2, mmu_idx=1)
10    at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:237
11 237         for (i = 0; i < levels; i++, ptshift -= ptidxbits) {
12 #查看循环参数:
13 (gdb) print i
14 $15 = -705808704 # 奇奇怪怪: 此时循环还没开始, i还未被赋初值, 也合理
15 (gdb) print ptshift
16 $16 = 18 # levels=3, (3-1)*9=18, 表示第一级偏移, 正确!
17 (gdb) print levels
```

```

18 $17 = 3 # levels=3, 确实是SV39三级页表, 正确!
19 (gdb) print ptidxbits
20 $18 = 9 # 每级9位索引, 正确!
21 (gdb) print /x base
22 $19 = 0x80206000 # 页表基地址

```

继续调试:

```

(gdb) until 238
get_physical_address (env=0x6434914bf040, physical=0x79cbd5ee21f0,
    prot=0x79cbd5ee21e4, addr=18446744072637907160, access_type=2,
    mmu_idx=1) at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:238
238         target_ulong idx = (addr >> (PGSHIFT + ptshift)) &
(gdb) print i
$20 = 0
(gdb) print /x idx
$21 = 0x643456b27ef9
(gdb) print /x (0xfffffffffc02000d8 >> (12 + 18)) & 0x1FF
$22 = 0x1ff
(gdb) until 242
get_physical_address (env=0x6434914bf040, physical=0x79cbd5ee21f0,
    prot=0x79cbd5ee21e4, addr=18446744072637907160, access_type=2,
    mmu_idx=1) at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:242
242         target_ulong pte_addr = base + idx * ptesize;
(gdb) print /x pte_addr
$23 = 0x1000
(gdb) print /x base + 511 * 8
$24 = 0x80206ff8

```

代码块

```

1 (gdb) until 238 #继续执行到计算 idx 的地方
2 get_physical_address (env=0x6434914bf040, physical=0x79cbd5ee21f0,
3     prot=0x79cbd5ee21e4, addr=18446744072637907160, access_type=2,
4     mmu_idx=1) at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:238
5 238         target_ulong idx = (addr >> (PGSHIFT + ptshift)) &
6 (gdb) print i
7 $20 = 0 #这次 i = 0, 正确!
8 (gdb) print /x idx
9 $21 = 0x643456b27ef9
10 (gdb) print /x (0xfffffffffc02000d8 >> (12 + 18)) & 0x1FF # 验证计算
11 $22 = 0x1ff
12 (gdb) until 242 #继续到计算 PTE 地址
13 get_physical_address (env=0x6434914bf040, physical=0x79cbd5ee21f0,
14     prot=0x79cbd5ee21e4, addr=18446744072637907160, access_type=2,
15     mmu_idx=1) at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:242
16 242         target_ulong pte_addr = base + idx * ptesize;
17 (gdb) print /x pte_addr
18 $23 = 0x1000
19 (gdb) print /x base + 511 * 8 #计算期望的 PTE 地址

```

```
20 $24 = 0x80206ff8
```

我们再确认一下当前的状态：

```
(gdb) print /x addr
$25 = 0xfffffffffc02000d8
(gdb) info locals
idx = 511
pte_addr = 4096
pte = 110175958270040
ppn = 110175958270480
mode = 1
base = 2149605376
levels = 3
ptidxbits = 9
ptesize = 8
vm = 8
sum = 0
mxr = 0
__func__ = "get_physical_address"
cs = 0x6434914b6630
va_bits = 39
mask = 67108863
masked_msbs = 67108863
ptshift = 18
i = 0
```

### 代码块

```
1 (gdb) print /x addr
2 $25 = 0xfffffffffc02000d8 #正确! ✓
3 (gdb) info locals
4 idx = 511 #VPN[2] = 511, 正确! ✓
5 pte_addr = 4096 #应该是base + idx * ptesize = 0x80206FF8, 有问题
6 pte = 110175958270040
7 ppn = 110175958270480
8 mode = 1
9 base = 2149605376
10 levels = 3
11 ptidxbits = 9
12 ptesize = 8
13 vm = 8
14 sum = 0
15 mxr = 0
16 __func__ = "get_physical_address"
17 cs = 0x6434914b6630
18 va_bits = 39
19 mask = 67108863
20 masked_msbs = 67108863
21 ptshift = 18
```

```
22 i = 0
```

现在 `pte_addr` 有问题，这说明GDB显示的 `pte_addr` 不是当前值，可能是变量还未被赋值，或GDB显示的是旧值，或实际内存地址不同。

我们继续努力：

```
(gdb) until 252
get_physical_address (env=0x6434914bf040, physical=0x79cbd5ee21f0,
    prot=0x79cbd5ee21e4, addr=18446744072637907160, access_type=2,
    mmu_idx=1) at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:252
252             target_ulong pte = ldq_phys(cs->as, pte_addr);
(gdb) print /x pte
$26 = 0x643456db0058
(gdb) print /x pte & 0x1
$27 = 0x0
(gdb) print /x base
$28 = 0x80206000
```

代码块

```
1 (gdb) until 252 #继续到读取PTE后
2 get_physical_address (env=0x6434914bf040, physical=0x79cbd5ee21f0,
3     prot=0x79cbd5ee21e4, addr=18446744072637907160, access_type=2,
4     mmu_idx=1) at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:252
5 252             target_ulong pte = ldq_phys(cs->as, pte_addr);
6 (gdb) print /x pte #查看PTE内容
7 $26 = 0x643456db0058
8 (gdb) print /x pte & 0x1 #检查是否有效
9 $27 = 0x0
10 (gdb) print /x base #如果是中间PTE, 查看base是否更新
11 $28 = 0x80206000
```

我们发现了关键问题：PTE无效……当前pte的值为0x643456db0058，V位为0，表示无效。R|W|X位：从值看可能都不为0，但因为V=0，所以无效。这就意味着：页表遍历失败，会触发页面错误，OS需要建立这个地址的页表映射。虽然PTE无效，但我们可以验证地址：

```
(gdb) print /x 0x80206000 + 511 * 8
$29 = 0x80206ff8
(gdb) x/1xg 0x80206ff8
0x80206ff8:      Cannot access memory at address 0x80206ff8
(gdb) x/1xg 0x80206000
0x80206000:      Cannot access memory at address 0x80206000
(gdb) x/1xg 0xfffffffffc02000d8
0xfffffffffc02000d8:      Cannot access memory at address 0xfffffffffc02000d8
```

代码块

```
1 (gdb) print /x 0x80206000 + 511 * 8 #查看PTE地址是否正确
```

```
2 $29 = 0x80206ff8
3 (gdb) x/1xg 0x80206ff8 #查看内存是否可访问
4 0x80206ff8: Cannot access memory at address 0x80206ff8
```

我们又发现了关键问题：内存访问失败。这说明物理地址 `0x80206ff8` 不存在或不可访问，QEMU 无法访问guest的物理内存，可能页表所在的物理内存区域未映射。我又试了其他几个之前出现过的地址，都是无法访问。

以下是大模型为我分析的原因：

1. 早期启动阶段：页表物理内存还没被正确映射到QEMU的地址空间
2. PMP限制：物理内存保护阻止了访问
3. QEMU内存管理：某些物理内存区域对host不可直接访问

到这里本来都想放弃了，但我注意到之前大模型给出的一句话”OS需要建立这个地址的页表映射“，那我们就让它去建立这个映射，然后记录调试过程就可以了呀！

向大模型求助后，它让我用finish命令从当前get\_physical\_address函数中跳出来：

```
(gdb) finish
Run till exit from #0  get_physical_address (env=0x6434914bf040,
    physical=0x79cbd5ee21f0, prot=0x79cbd5ee21e4,
    addr=18446744072637907160, access_type=2, mmu_idx=1)
  at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:252
0x00006434568d4c9e in riscv_cpu_tlb_fill (cs=0x6434914b6630,
    address=18446744072637907160, size=0, access_type=MMU_INST_FETCH,
    mmu_idx=1, probe=false, retaddr=0)
  at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:451
451      ret = get_physical_address(env, &pa, &prot, address, access_type
, mmu_idx);
Value returned is $30 = 0
```

#### 代码块

```
1 (gdb) finish
2 Run till exit from #0  get_physical_address (env=0x6434914bf040,
3     physical=0x79cbd5ee21f0, prot=0x79cbd5ee21e4,
4     addr=18446744072637907160, access_type=2, mmu_idx=1)
5     at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:252
6 0x00006434568d4c9e in riscv_cpu_tlb_fill (cs=0x6434914b6630,
7     address=18446744072637907160, size=0, access_type=MMU_INST_FETCH,
8     mmu_idx=1, probe=false, retaddr=0)
9     at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:451
10 451      ret = get_physical_address(env, &pa, &prot, address, access_type,
11     mmu_idx);
11 Value returned is $30 = 0 #返回结果为0，表示还是成功了
```

通过执行 `finish` 命令，我们从正在单步调试的 `get_physical_address` 函数中返回到其调用者 `riscv_cpu_tlb_fill`。输出显示：

1. `get_physical_address` 函数返回值为 `0` (`$30 = 0`)
2. 这个返回值对应源码中的 `TRANSLATE_SUCCESS`，表示页表遍历成功
3. 返回到 `riscv_cpu_tlb_fill` 函数的第451行，这是调用 `get_physical_address` 之后的位置
4. 传入的参数包括：虚拟地址 `18446744072637907160` (即 `0xfffffffffc02000d8`)、访问类型 `MMU_INST_FETCH` (值为2，表示指令获取)、MMU索引 `1` (S模式)
5. 返回值 `0` 将被存储在局部变量 `ret` 中，后续代码会根据这个值决定是建立TLB映射还是触发异常

这表明尽管之前观察到的页表项 (PTE) 的V位为0，但在实际的页表遍历过程中，该地址的翻译最终成功完成。可能是由于多级页表遍历找到了有效的最终页表项，或者在并发环境下页表被其他CPU核心更新。

然后我输入bt命令，查看当前的调用栈：

```
(gdb) bt
#0 0x00006434568d4c9e in riscv_cpu_tlb_fill (cs=0x6434914b6630,
    address=18446744072637907160, size=0, access_type=MMU_INST_FETCH,
    mmu_idx=1, probe=false, retaddr=0)
    at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:451
#1 0x000064345681afdf in tlb_fill (cpu=0x6434914b6630,
    addr=18446744072637907160, size=0, access_type=MMU_INST_FETCH,
    mmu_idx=1, retaddr=0)
    at /mnt/d/For_OS/qemu-4.1.1/accel/tcg/cputlb.c:878
#2 0x000064345681b689 in get_page_addr_code (env=0x6434914bf040,
    addr=18446744072637907160)
    at /mnt/d/For_OS/qemu-4.1.1/accel/tcg/cputlb.c:1032
#3 0x00006434568400ef in tb_htable_lookup (cpu=0x6434914b6630,
    pc=18446744072637907160, cs_base=0, flags=24577, cf_mask=4278714368)
    at /mnt/d/For_OS/qemu-4.1.1/accel/tcg/cpu-exec.c:339
#4 0x00006434568390c5 in tb_lookup_cpu_state (cpu=0x6434914b6630,
    pc=0x79cbd5ee23e0, cs_base=0x79cbd5ee23d8, flags=0x79cbd5ee23d4,
    cf_mask=4278714368)
    at /mnt/d/For_OS/qemu-4.1.1/include/exec/tb-lookup.h:43
#5 0x000064345683947e in helper_lookup_tb_ptr (env=0x6434914bf040)
    at /mnt/d/For_OS/qemu-4.1.1/accel/tcg/tcg-runtime.c:154
#6 0x000079cbce050df7 in code_gen_buffer ()
#7 0x000064345683fb99 in cpu_tb_exec (cpu=0x6434914b6630,
    itb=0x79cbce050d00 <code_gen_buffer+330963>)
    at /mnt/d/For_OS/qemu-4.1.1/accel/tcg/cpu-exec.c:173
#8 0x00006434568409df in cpu_loop_exec_tb (cpu=0x6434914b6630,
    tb=0x79cbce050d00 <code_gen_buffer+330963>, last_tb=0x79cbd5ee2988,
    tb_exit=0x79cbd5ee2980)
    at /mnt/d/For_OS/qemu-4.1.1/accel/tcg/cpu-exec.c:621
#9 0x0000643456840d05 in cpu_exec (cpu=0x6434914b6630)
    at /mnt/d/For_OS/qemu-4.1.1/accel/tcg/cpu-exec.c:732
#10 0x00006434567f2f3f in tcg_cpu_exec (cpu=0x6434914b6630)
    at /mnt/d/For_OS/qemu-4.1.1/cpus.c:1435
#11 0x00006434567f37f8 in qemu_tcg_cpu_thread_fn (arg=0x6434914b6630)
    at /mnt/d/For_OS/qemu-4.1.1/cpus.c:1743
#12 0x0000643456c768d7 in qemu_thread_start (args=0x6434914cccc0)
    at util/qemu-thread-posix.c:502
#13 0x000079cbd689caa4 in start_thread (arg=<optimized out>)
    at ./nptl/pthread_create.c:447
#14 0x000079cbd6929c6c in clone3 ()
    at ../sysdeps/unix/sysv/linux/x86_64/clone3.S:78
```

## 代码块

```
1 (gdb) bt
2 #0 0x00006434568d4c9e in riscv_cpu_tlb_fill (cs=0x6434914b6630,
3     address=18446744072637907160, size=0, access_type=MMU_INST_FETCH,
4     mmu_idx=1, probe=false, retaddr=0)
5     at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:451
6 #1 0x000064345681afdf in tlb_fill (cpu=0x6434914b6630,
7     addr=18446744072637907160, size=0, access_type=MMU_INST_FETCH,
8     mmu_idx=1, retaddr=0)
9     at /mnt/d/For_OS/qemu-4.1.1/accel/tcg/cputlb.c:878
```

```

10 #2 0x000064345681b689 in get_page_addr_code (env=0x6434914bf040,
11     addr=18446744072637907160)
12     at /mnt/d/For_OS/qemu-4.1.1/accel/tcg/cputlb.c:1032
13 #3 0x00006434568400ef in tb_htable_lookup (cpu=0x6434914b6630,
14     pc=18446744072637907160, cs_base=0, flags=24577, cf_mask=4278714368)
15     at /mnt/d/For_OS/qemu-4.1.1/accel/tcg/cpu-exec.c:339
16 #4 0x00006434568390c5 in tb_lookup_cpu_state (cpu=0x6434914b6630,
17     pc=0x79cbd5ee23e0, cs_base=0x79cbd5ee23d8, flags=0x79cbd5ee23d4,
18     cf_mask=4278714368)
19     at /mnt/d/For_OS/qemu-4.1.1/include/exec/tb-lookup.h:43
20 #5 0x000064345683947e in helper_lookup_tb_ptr (env=0x6434914bf040)
21     at /mnt/d/For_OS/qemu-4.1.1/accel/tcg/tcg-runtime.c:154
22 #6 0x000079cbce050df7 in code_gen_buffer ()
23 #7 0x000064345683fb99 in cpu_tb_exec (cpu=0x6434914b6630,
24     itb=0x79cbce050d00 <code_gen_buffer+330963>)
25     at /mnt/d/For_OS/qemu-4.1.1/accel/tcg/cpu-exec.c:173
26 #8 0x00006434568409df in cpu_loop_exec_tb (cpu=0x6434914b6630,
27     tb=0x79cbce050d00 <code_gen_buffer+330963>, last_tb=0x79cbd5ee2988,
28     tb_exit=0x79cbd5ee2980)
29     at /mnt/d/For_OS/qemu-4.1.1/accel/tcg/cpu-exec.c:621
30 #9 0x0000643456840d05 in cpu_exec (cpu=0x6434914b6630)
31     at /mnt/d/For_OS/qemu-4.1.1/accel/tcg/cpu-exec.c:732
32 #10 0x00006434567f2f3f in tcg_cpu_exec (cpu=0x6434914b6630)
33     at /mnt/d/For_OS/qemu-4.1.1/cpus.c:1435
34 #11 0x00006434567f37f8 in qemu_tcg_cpu_thread_fn (arg=0x6434914b6630)
35     at /mnt/d/For_OS/qemu-4.1.1/cpus.c:1743
36 #12 0x0000643456c768d7 in qemu_thread_start (args=0x6434914cccc0)
37     at util/qemu-thread-posix.c:502
38 #13 0x000079cbd689caa4 in start_thread (arg=<optimized out>)
39     at ./nptl/pthread_create.c:447
40 #14 0x000079cbd6929c6c in clone3 ()
41     at ../sysdeps/unix/sysv/linux/x86_64/clone3.S:78

```

对于上述调用栈：

**底层系统调用层 (#13-#14) :**

- QEMU用户态线程的创建和调度 (`start_thread`, `clone3`)

**QEMU线程调度层 (#10-#12) :**

- QEMU的TCG (Tiny Code Generator) CPU执行循环
- `qemu_tcg_cpu_thread_fn` : QEMU的TCG CPU线程主函数
- `tcg_cpu_exec` : TCG模式的CPU执行入口

**CPU执行核心层 (#7-#9) :**

- `cpu_exec` : CPU主执行循环

- `cpu_loop_exec_tb`：执行翻译块（Translation Block）
- `cpu_tb_exec`：执行单个翻译块

### 代码生成与查找层 (#3-#6) :

- `helper_lookup_tb_ptr`：帮助函数查找翻译块指针
- `tb_lookup_cpu_state`：根据CPU状态查找翻译块
- `tb_htable_lookup`：在翻译块哈希表中查找
- `code_gen_buffer`：QEMU动态生成的机器代码区域

### 内存管理/TLB层 (#0-#2) : (核心)

- `get_page_addr_code`：获取代码页的物理地址
- `tlb_fill`：TLB填充函数，处理TLB缺失
- `riscv_cpu_tlb_fill`：RISC-V架构特定的TLB填充实现（当前执行位置）

### 上述信息说明：

1. 当前正执行在 `riscv_cpu_tlb_fill` 中（栈帧#0），这是处理虚拟地址翻译的入口
2. 触发原因是获取代码页地址（`get_page_addr_code`），说明是指令获取导致的页表访问
3. 虚拟地址为 `0xfffffffffc02000d8`，位于内核空间
4. 调用链展示了QEMU从执行用户代码→翻译块管理→内存访问→页表遍历的完整路径
5. 这表明操作系统内核代码执行时遇到了页表缺失，需要建立映射

---

接下来我输入多个next（next会执行当前代码，遇到函数时会执行整个函数，并在函数返回后暂停）：

```

(gdb) next
453         if (mode == PRV_M && access_type != MMU_INST_FETCH) {
(gdb) next
459             qemu_log_mask(CPU_LOG_MMU,
(gdb) next
463                 if (riscv_feature(env, RISCV_FEATURE_PMP) &&
(gdb) next
465                     !pmp_hart_has_privs(env, pa, size, 1 << access_type, mode))
{
(gdb) next
464             (ret == TRANSLATE_SUCCESS) &&
(gdb) next
468                 if (ret == TRANSLATE_PMP_FAIL) {
(gdb) next
471                     if (ret == TRANSLATE_SUCCESS) {
(gdb) next
472                         tlb_set_page(cs, address & TARGET_PAGE_MASK, pa & TARGET_PAG
E_MASK,
(gdb) next
474                     return true;
(gdb) next
495     }
(gdb) next
500     tlb_fill (cpu=0x6434914b6630, addr=18446744072637907160, size=0,
    access_type=MMU_INST_FETCH, mmu_idx=1, retaddr=0)
    at /mnt/d/For_OS/qemu-4.1.1/accel/tcg/cputlb.c:879
879         assert(ok);
(gdb) c
Continuing.

```

```

Thread 2 "qemu-system-ris" hit Breakpoint 4, get_physical_address (
  env=0x6434914bf040, physical=0x79cbd5ee2290, prot=0x79cbd5ee2284,
  addr=18446744072637931512, access_type=1, mmu_idx=1)
  at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:237
237         for (i = 0; i < levels; i++, ptshift -= ptidxbits) {

```

## 代码块

```

1  (gdb) next
2  453         if (mode == PRV_M && access_type != MMU_INST_FETCH) {
3  (gdb) next
4  459             qemu_log_mask(CPU_LOG_MMU,
5  (gdb) next
6  463                 if (riscv_feature(env, RISCV_FEATURE_PMP) &&
7  (gdb) next
8  465                     !pmp_hart_has_privs(env, pa, size, 1 << access_type, mode)) {
9  (gdb) next
10 464             (ret == TRANSLATE_SUCCESS) &&
11 (gdb) next
12 468                 if (ret == TRANSLATE_PMP_FAIL) {
13 (gdb) next
14 471                     if (ret == TRANSLATE_SUCCESS) {
15 (gdb) next

```

```

16 472             tlb_set_page(cs, address & TARGET_PAGE_MASK, pa &
17 TARGET_PAGE_MASK,
18 474             return true;
19 (gdb) next
20 495 }
21 (gdb) next
22 tlb_fill (cpu=0x6434914b6630, addr=18446744072637907160, size=0,
23     access_type=MMU_INST_FETCH, mmu_idx=1, retaddr=0)
24     at /mnt/d/For_OS/qemu-4.1.1/accel/tcg/cputlb.c:879
25 879         assert(ok); #TLB填充成功
26 (gdb) c
27 Continuing.
28
29 Thread 2 "qemu-system-riscv" hit Breakpoint 4, get_physical_address (
30     env=0x6434914bf040, physical=0x79cbd5ee2290, prot=0x79cbd5ee2284,
31     addr=18446744072637931512, access_type=1, mmu_idx=1)
32     at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:237
33 237         for (i = 0; i < levels; i++, ptshift -= ptidxbits) {

```

发现一直 `next` 下去，最终到了 `tlb_fill` 函数。这是因为 `riscv_cpu_tlb_fill` 返回后，控制流返回到它的调用者 `tlb_fill` 处。调用关系是：

#### 代码块

```

1 tlb_fill() # 通用TLB填充函数
2   ↓
3 riscv_cpu_tlb_fill() # RISC-V架构特定的实现
4   ↓ (返回true)
5 tlb_fill() # 继续执行879行

```

其中879行的 `assert(ok);` 表示TLB填充成功。接下来我们输入c继续执行，发现很快就遇到了新的内存访问（虚拟地址 `0xfffffffffc02000f8`），触发了新的页表遍历，命中断点4，再次进入了 `get_physical_address`。

然后我们像之前一样继续去调试：

```

(gdb) next
238          target_ulong idx = (addr >> (PGSHIFT + ptshift)) &
(gdb) print i
$35 = 0
(gdb) print /x addr
$36 = 0xfffffffffc0205ff8
(gdb) print ptshift
$37 = 18
(gdb) print /x (addr >> (12 + 18)) & 0x1FF
$38 = 0x1ff
(gdb) until 252
get_physical_address (env=0x6434914bf040, physical=0x79cbd5ee2290,
    prot=0x79cbd5ee2284, addr=18446744072637931512, access_type=1,
    mmu_idx=1) at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:252
252          target_ulong pte = ldq_phys(cs->as, pte_addr);
(gdb) print /x base
$39 = 0x80206000
(gdb) print /x idx
$40 = 0x1ff
(gdb) print /x pte_addr
$41 = 0x80206ff8
(gdb) next
254          target_ulong ppn = pte >> PTE_PPN_SHIFT;
(gdb) print /x pte
$42 = 0x200000cf
(gdb) print /x pte & 0xffff
$43 = 0xcf
(gdb) print /x pte & 1
$44 = 0x1
(gdb) print /x pte & 0xf
$45 = 0xf

```

## 代码块

```

1  (gdb) next #计算页表索引 idx: idx = (addr >> (12 + 18)) & 0x1FF
2  238          target_ulong idx = (addr >> (PGSHIFT + ptshift)) &
3  (gdb) print i #验证参数
4  $35 = 0      #i = 0, 表示在第0级页表
5  (gdb) print /x addr
6  $36 = 0xfffffffffc0205ff8 #虚拟地址
7  (gdb) print ptshift
8  $37 = 18    #当前偏移量
9  (gdb) print /x (addr >> (12 + 18)) & 0x1FF
10 $38 = 0x1ff #计算得到索引511
11 (gdb) until 252 #计算PTE地址
12 get_physical_address (env=0x6434914bf040, physical=0x79cbd5ee2290,
13     prot=0x79cbd5ee2284, addr=18446744072637931512, access_type=1,
14     mmu_idx=1) at /mnt/d/For_OS/qemu-4.1.1/target/riscv/cpu_helper.c:252
15 252          target_ulong pte = ldq_phys(cs->as, pte_addr);
16 (gdb) print /x base
17 $39 = 0x80206000      #页表基址
18 (gdb) print /x idx

```

```

19 $40 = 0x1ff          #页表索引
20 (gdb) print /x pte_addr
21 $41 = 0x80206ff8      #PTE物理地址
22 # 计算: 0x80206000 + 511 * 8 = 0x80206000 + 0xFF8 = 0x80206ff8
23 (gdb) next
24 254                  target_ulong ppn = pte >> PTE_PPN_SHIFT;
25 (gdb) print /x pte
26 $42 = 0x200000cf    #页表项值
27 (gdb) print /x pte & 0xffff
28 $43 = 0xcf          #低12位标志位
29 (gdb) print /x pte & 1
30 $44 = 0x1            #V位=1, 有效!
31 (gdb) print /x pte & 0xf
32 $45 = 0xf            #低4位: V=1, R=0, W=1, X=1

```

上述的步骤：计算页表索引→验证参数→计算PTE地址→读取PTE并分析。

PTE标志位详细分析：

#### 代码块

```

1 0xcf = 1100 11112
2 位分解:
3 - 位0 (0x01): V = 1    有效
4 - 位1 (0x02): R = 0    不可读
5 - 位2 (0x04): W = 1    可写
6 - 位3 (0x08): X = 1    可执行
7 - 位4 (0x10): U = 0    超级用户
8 - 位5 (0x20): G = 0    非全局
9 - 位6 (0x40): A = 1    已访问
10 - 位7 (0x80): D = 1   已修改

```

为什么之前PTE的V位是0，现在是1了：

#### 代码块

```

1 时间点1: 访问 0xfffffffffc02000d8 (指令获取)
2     → PTE.V = 0 (无效)
3     → 触发页错误异常 (Instruction Page Fault)
4
5 时间点2: 访问 0xfffffffffc0205ff8 (数据加载)
6     → PTE.V = 1 (有效)
7     → 页表遍历成功

```

大模型说这个正常OS的行为：

1. 按需分配：页面在第一次访问时才分配
  2. 惰性加载：页表项初始化为无效，触发异常时才建立映射
  3. 写时复制：某些页面初始为只读，写时触发异常再设为可写
  4. 内存共享：通过精心控制页表权限实现
- 

然后我又执行了很多步的next：

```
(gdb) next
256          if (!(pte & PTE_V)) {
(gdb) next
259          } else if (!(pte & (PTE_R | PTE_W | PTE_X))) {
(gdb) print /x pte & (PTE_R | PTE_W | PTE_X)
No symbol "PTE_R" in current context.
(gdb) print /x (pte & 0xffff) & (0x002 | 0x004 | 0x008)
$46 = 0xe
(gdb) next
262          } else if ((pte & (PTE_R | PTE_W | PTE_X)) == PTE_W) {
(gdb) next
265          } else if ((pte & (PTE_R | PTE_W | PTE_X)) == (PTE_W | PTE_X)) {
(gdb) next
268          } else if ((pte & PTE_U) && ((mode != PRV_U) &&
(gdb) next
273          } else if (!(pte & PTE_U) && (mode != PRV_S)) {
(gdb) next
276          } else if (ppn & ((1ULL << ptshift) - 1)) {
(gdb) next
279          } else if (access_type == MMU_DATA_LOAD && !(pte & PTE_R) ||
(gdb) next
283          } else if (access_type == MMU_DATA_STORE && !(pte & PTE_W)) {
(gdb) next
286          } else if (access_type == MMU_INST_FETCH && !(pte & PTE_X)) {
(gdb) next
292          (access_type == MMU_DATA_STORE ? PTE_D : 0);
(gdb) next
291          target_ulong updated_pte = pte | PTE_A |
(gdb) next
295          if (updated_pte != pte) {
333          target_ulong vpn = addr >> PGSHIFT;
334          *physical = (ppn | (vpn & ((1L << ptshift) - 1))) << PGSHIFT;
337          if ((pte & PTE_R) || ((pte & PTE_X) && mxr)) {
(gdb) next
338              *prot |= PAGE_READ;
340          if ((pte & PTE_X)) {
341              *prot |= PAGE_EXEC;
345          if ((pte & PTE_W) &&
347              *prot |= PAGE_WRITE;
(gdb) next
349          return TRANSLATE_SUCCESS;
```

可以看到，这次程序跳过了所有失败分支，执行了成功路径，最终返回的是 `return TRANSLATE_SUCCESS;` 说明这回的虚拟地址到物理地址的翻译成功了。

我们看一下现在的参数值：

```
(gdb) info locals
updated_pte = 536871119
vpn = 4503599627108869
idx = 511
pte_addr = 2149609464
pte = 536871119
ppn = 524288
mode = 1
base = 2149605376
levels = 3
ptidxbits = 9
ptesize = 8
vm = 8
sum = 0
mxr = 0
__func__ = "get_physical_address"
cs = 0x6434914b6630
va_bits = 39
mask = 67108863
masked_msbs = 67108863
ptshift = 18
i = 0
(gdb) print /x *physical
$47 = 0x80205000
(gdb) print /x *prot
$48 = 0x7
(gdb) print /x addr
$49 = 0xfffffffffc0205ff8
(gdb) print /x ppn
$50 = 0x80000
(gdb) print ptshift
$51 = 18
(gdb) print /x vpn
$52 = 0xfffffffffc0205
(gdb) print /x (1L << ptshift) - 1
$53 = 0x3ffff
(gdb) print /x ppn | (vpn & ((1L << ptshift) - 1))
$54 = 0x80205
(gdb) print /x (ppn | (vpn & ((1L << ptshift) - 1))) << 12
$55 = 0x80205000
(gdb) print /x addr >> 12
$56 = 0xfffffffffc0205
(gdb) print /x vpn & 0x3ffff
$57 = 0x205
```

## 代码块

```
1 (gdb) info locals
2 updated_pte = 536871119
3 vpn = 4503599627108869
4 idx = 511
5 pte_addr = 2149609464
6 pte = 536871119
7 ppn = 524288
8 mode = 1
9 base = 2149605376
10 levels = 3
```

```
11 ptidxbits = 9
12 ptesize = 8
13 vm = 8
14 sum = 0
15 mxr = 0
16 __func__ = "get_physical_address"
17 cs = 0x6434914b6630
18 va_bits = 39
19 mask = 67108863
20 masked_msbs = 67108863
21 ptshift = 18
22 i = 0
23 (gdb) print /x *physical
24 $47 = 0x80205000
25 (gdb) print /x *prot
26 $48 = 0x7
27 (gdb) print /x addr
28 $49 = 0xfffffffffc0205ff8
29 (gdb) print /x ppn
30 $50 = 0x80000
31 (gdb) print ptshift
32 $51 = 18
33 (gdb) print /x vpn
34 $52 = 0xfffffffffc0205
35 (gdb) print /x (1L << ptshift) - 1
36 $53 = 0x3fff
37 (gdb) print /x ppn | (vpn & ((1L << ptshift) - 1))
38 $54 = 0x80205
39 (gdb) print /x (ppn | (vpn & ((1L << ptshift) - 1))) << 12
40 $55 = 0x80205000
41 (gdb) print /x addr >> 12
42 $56 = 0xfffffffffc0205
43 (gdb) print /x vpn & 0x3fff
44 $57 = 0x205
```

## 核心参数验证：

参数	值 (十六进制)	解释	正确性
虚拟地址	0xfffffffffc0205ff8	要翻译的地址	✓
PTE值	0x200000cf (536871119)	页表项	✓
PPN	0x80000 (524288)	物理页号	✓ (从PTE提取)
物理地址	0x80205000	计算结果	✓
权限	0x7	读+写+执行	✓

### 代码块

```

1 PTE = 0x200000cf (PTE是页表项)
2 PPN = (pte >> 10) & 0xffffffffffff = 0x80000 # ✓ (PPN是物理页号)
3 VPN偏移 = (addr >> 12) & 0xffff = 0x205 (多级页表遍历过程中，当前级别页表索引未覆盖
    的那部分虚拟地址位)
4 物理页号 = PPN | VPN偏移 = 0x80000 | 0x205 = 0x80205
5 物理地址 = 物理页号 << 12 = 0x80205000 # ✓ 与输出一致!

```

以上我们就成功跟踪了本次虚拟地址到物理地址的转换！

我们总结一下关键调试步骤：

- 触发了页表访问：虚拟地址 `0xfffffffffc0205ff8` 的访问
- 进入页表遍历：`get_physical_address` 函数
- 读取页表项：从物理地址 `0x80206ff8` 读取 PTE `0x200000cf`
- 分析PTE标志：
  - V=1 (有效)
  - R=0, W=1, X=1 (可写、可执行但不可读？)
  - U=0 (超级用户)
  - A=1, D=1 (已访问、已修改)
- 通过权限检查：尽管看起来R=0，但可能由于 `mrx` 或其他机制通过
- 计算物理地址：
  - VPN = `0xfffffffffc0205`
  - PPN = `0x80000`

- 物理地址 = `0x80205000`

## 7. 建立TLB映射：权限为 `0x7` (读+写+执行)

更重要的是，发现了一个关键信息：这是一个1GB大页映射！

- 在第一级页表 ( $i=0$ ,  $ptshift=18$ ) 就找到了叶子PTE
- 虚拟地址范围：`0xfffffffffc0000000` 到 `0xffffffffffffffff` → 物理地址范围：`0x80000000` 到 `0xbfffffff`
- 这是操作系统内核空间的典型映射方式

调试成果：

- 理解了RISC-V三级页表遍历过程
  - 观察了PTE标志位的检查和解析
  - 验证了虚拟地址到物理地址的计算公式
  - 发现了操作系统使用1GB大页优化内核映射
  - 跟踪了TLB缓存的建立过程
- 

## 指导书问题

### 3. 是否能够在qemu-4.1.1的源码中找到模拟CPU查找TLB的C代码？

在 `qemu-4.1.1/accel/tcg/cputlb.c` 中

#### 主要 TLB 查找相关函数

1. `tbl_set_page_with_attrs()`
  - 功能：向 TLB 中添加一个新的条目（虚拟地址 → 物理地址映射）。
  - 在 TCG 生成代码时调用，用于填充 TLB。
  - 包含地址映射、权限检查、缓存设置等逻辑。
2. `tbl_fill()`
  - 功能：当 TLB 未命中时，触发异常处理并填充 TLB。
  - 调用 CPU 特定的 `tbl_fill` 方法（例如 `cpu_arm_tlb_fill`）。
3. `get_page_addr_code()`
  - 功能：获取代码页的物理地址，用于指令翻译。
  - 如果 TLB 命中，返回物理地址；否则触发 TLB 填充。
4. `probe_write()`

- 功能：检查写访问权限，如果未命中则触发 TLB 填充。

#### 5. `tlb_vaddr_to_host()`

- 功能：将虚拟地址转换为宿主（host）地址指针。
- 如果 TLB 命中且是 RAM 访问，返回 host 指针；否则返回 `NULL`。

#### 6. `atomic_mmu_lookup()`

- 功能：原子操作（RMW）的 TLB 查找，支持对齐检查和权限验证。

#### 7. `load_helper()` / `store_helper()`

- 功能：通用的加载和存储辅助函数，封装了 TLB 查找、对齐检查、IO 访问、未对齐访问等逻辑。

## ✓ TLB 数据结构

在文件中涉及以下关键数据结构：

- `CPUTLBEntry`：表示一个 TLB 条目，包含地址、权限等信息。
  - `CPUTLBDesc`：描述 TLB 的状态（如大小、使用条目数等）。
  - `CPUIOTLBEntry`：IO 地址映射条目。
- 

## 4. 仍然是 TLB，qemu 中模拟出来的 TLB 和我们真实 CPU 中的 TLB 有什么逻辑上的区别？

对比维度	QEMU 模拟 TLB	真实 CPU TLB	核心差异说明
根本性质	软件实现的缓存机制	硬件实现的缓存单元	QEMU TLB是软件算法，真实TLB是物理电路
设计目的	加速虚拟机的地址翻译	加速真实程序的地址翻译	QEMU是为仿真加速，CPU是为应用加速
翻译层级	两级翻译： Guest虚拟地址 → Guest物理地址 → Host虚拟地址	一级翻译： CPU虚拟地址 → 物理地址	QEMU需要多一层Guest到Host的映射
物理位置	在QEMU进程的虚拟地址空间中	在CPU芯片内部（专用SRAM）	一个在内存，一个在芯片
访问速度	纳秒级（内存访问）	皮秒级（片上缓存访问）	硬件TLB快2-3个数量级
大小管理	动态可调整（ <code>tlb_mmu_resize_locked</code> ）	固定大小（硬件设计时确定）	QEMU可根据负载调整，硬件TLB固定
失效机制	软件显式调用 <code>tlb_flush*</code> ()函数	硬件自动管理+软件指令（如INVLPG）	QEMU完全软件控制，真实CPU软硬结合
命中处理	返回Host内存指针，继续软件执行	直接提供物理地址给ALU	QEMU仍需软件处理，CPU硬件直接使用
未命中处理	调用 <code>tlb_fill</code> ()触发软件异常处理	触发Page Walk硬件单元遍历页表	QEMU走异常处理流程，CPU走硬件流程
替换策略	软件实现（可自定义算法）	硬件固定（通常为LRU或随机）	QEMU灵活可调，CPU策略固化
与MMU关系	模拟Guest的MMU行为	是CPU MMU的一部分	QEMU TLB模拟MMU功能，真实TLB是MMU组件
异常触发	可能导致QEMU内部的异常处理	可能导致CPU异常（如Page Fault）	QEMU异常在软件层，CPU异常在硬件层